



HIGH PERFORMANCE AND PARALLEL COMPUTING REPORT

CAB401: High Performance and Parallel Computing

OCTOBER 27, 2017
MICHAEL CARTWRIGHT
N8869871

Table of Contents

Paragraph	Page No.
1 Introduction	1
1.1 Software Architecture.....	1
1.1.1 Time	1
1.1.2 Initialise	1
1.1.3 Input Filename and Read Input File	1
1.1.4 Compare All Bacteria	2
1.1.5 Bacteria Class	3
1.1.6 Compare Bacteria	5
2 Bioinformatics – Genome Similarity Using Frequency Vectors	7
2.1 Software, Compiler, Hardware and Techniques.....	7
2.1.1 Software	7
2.1.2 Compiler.....	7
2.1.3 Hardware	7
2.1.4 Techniques	7
2.2 Analysis of Code and Parallelisation Process	8
2.2.1 Successful Parallelisation	8
2.2.2 Further Analysed Parallelisation Attempts	9
2.2.2.1 Two Functions at the Same Time in Main.....	9
2.2.2.2 Parallelisation Bacteria Class	9
2.2.2.3 Parallelisation Comparing Two Bacteria	10
3 Code Implemented.....	11
3.1 Code Changes in Compare All Bacteria Function.....	11
4 Results	15
4.1 Time Taken for Execution of Sequential and Parallelised Applications.....	15
4.2 Profiling Reports	15
4.3 Speed Up Graph	17
5 Reflection	18
5.1 Issues Encountered	18
6 Appendix	20
6.1 Original Code	20
6.2 High Performance and Parallel Code.....	25

6.3	Detailed Software Architecture.....	30
6.4	Bacteria Class – Part 1 Software Architecture	31
6.5	Bacteria Class – Part 2 Software Architecture	32
6.6	Compare Bacteria Software Architecture	33

List of Figures

Figure 1: Basic View of Software Architecture	1
Figure 2: Read Input File Software Architecture	2
Figure 3: Compare All Bacteria Software Architecture	2
Figure 4: AcMNPV Animo Acid File Content Example	3
Figure 5: Get Data from *.faa File	3
Figure 6: Calculate Variables in Bacteria Class – Part 1	4
Figure 7: Other Values Calculated in Bacteria Class	4
Figure 8: Bacteria Class Ending.....	5
Figure 9: Detail View of Compare Bacteria	6
Figure 10: Total CPU (ms, %) of Functions in Bioinformatics – Genome Similarity Using Frequency Vectors Application	6
Figure 11: Original Code for Compare All Bacteria Function	11
Figure 12: Parallelised Code for Compare All Bacteria Function	11
Figure 13: Original Loading Bacteria.....	12
Figure 14: Parallel Loading Bacteria Four Threads Max	12
Figure 15: Parallel Loading Bacteria Eight Threads Max	12
Figure 16: Sequential Application Memory and CPU Usage	13
Figure 17: Parallelised Application Memory and CPU Usage	13
Figure 18: Parallelised Nested For-Loop in Order.....	14
Figure 19: Correlation in Order.....	14
Figure 20: Correlation not in Order.....	14
Figure 21: Sequential CPU Profiler	16

Figure 22: Parallelised CPU Profiler.....	16
Figure 23: Speed Up Graph of Sequential vs Parallelised Over Threads	17
Figure 24: Restructured Nested For-Loop Attempt.....	18

List of Tables

Table 1: Sequential and Different Thread Count for Parallelised Application

Run-Time	15
----------------	----

1 Introduction

The application chosen to be parallelised was the Bioinformatics – Genome Similarity Using Frequency Vectors. This application’s purpose is to compare different types of bacteria to each other by calculating the correlation between two different bacteria.

1.1 Software Architecture

Figure 1 below shows the basic process of Bioinformatics – Genome Similarity Using Frequency Vectors’ architecture. The detailed software architecture can be found in full in section 6.3 of the Appendix of this report.

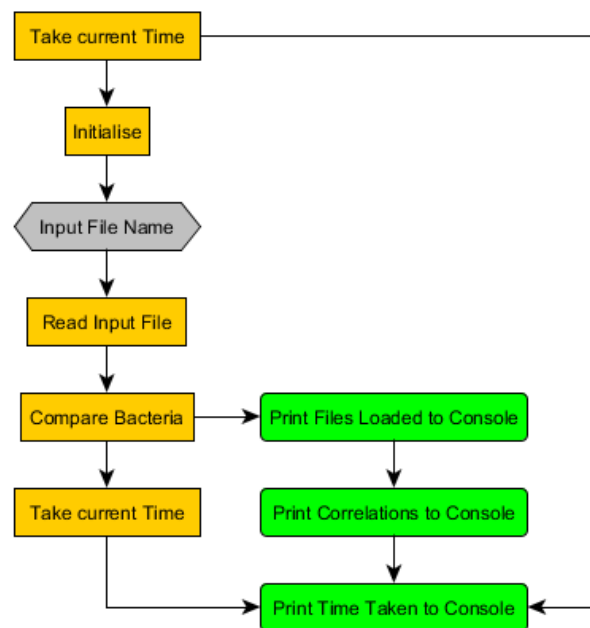


Figure 1: Basic View of Software Architecture

1.1.1 Time

The application takes the time before and after the application and the difference is calculated. This is printed to the console to notify the user how long the program took to complete the Bioinformatics – Genome Similarity Using Frequency Vectors application.

1.1.2 Initialise

This method initialises long variables M1 and M, based on the predefined value for AA_NUMBER and LEN, for the class Bacteria.

1.1.3 Input Filename and Read Input File

The program takes a users’ input argument. This argument is the filename for a text file. This text file contains the number of bacteria and the names of the associated bacteria files. The first line contains the number of bacteria. From the second line down, these lines contain the name of each *.faa file. The extension *.faa is used by the NCBI for FASTA animo acids. In this application, forty-one protein FASTA files are being used. Figure 2, on the next page, shows the detail view software architecture for this function.

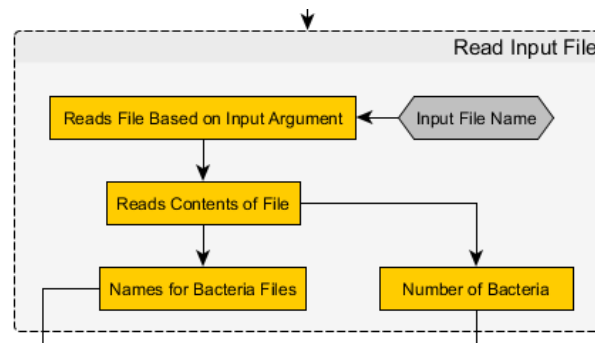
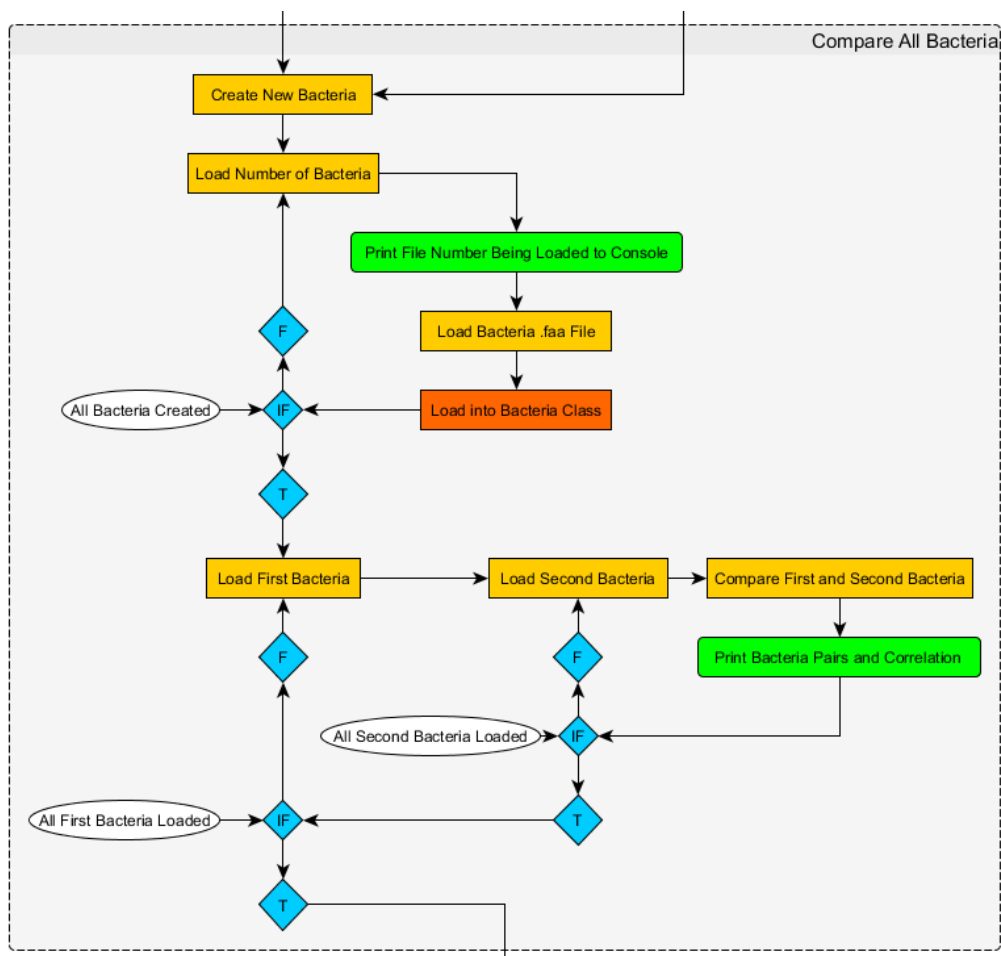


Figure 2: Read Input File Software Architecture

1.1.4 Compare All Bacteria

This component of the application is where most of the runtime and computational work is located. This function does three things. The first creates the number of classes of bacteria based on the number of bacteria to be read. The second component is made of two parts. The first part is printing to the console which file is being loaded. The second part is loading the *.faa files and loading the data into their respective bacteria class. The third calls for the compare bacteria function which calculates the correlation between two pairs of different bacteria and then printing the result to the console. Figure 3 on the next page shows a detailed view of the CompareAllBacteria function.



1.1.5 Bacteria Class

The bacteria files are opened and vectors are initialised. The bacteria class holds data based on the contents in the *.faa files. The class obtains the necessary characters from the respective *.faa file. Figure 4, shown below, shows some of the contents of file AcMNPV.faa.

```

1 >gi|9627743|ref|NP_054030.1| protein tyrosine phosphatase [Autographa californica nucleopolyhedrovirus]
2 MFPARWHNYLQCGQVIKDSNLICFKITPLRPELFAYVTSEEDVMTAEQIVKQNPISIGAIIDLNTSKYYDG
3 VHFLRAGLLYKKIQVPGQTLPPESIVQEFIDTVKEFTEKCPGMLVGVHCTHGINRTGYMVCRYLMHTLGI
4 APQEAIDRFKARGHKIERQNYVQDLLI
5 >gi|9627744|ref|NP_054031.1| baculovirus repeated ORF [Autographa californica nucleopolyhedrovirus]
6 MARVKIGEFKFGEDTFNLRYVLERDQQVRFVAKDVANSLKYTVCDKAIRVHVDNKKYKSLFEQTIQNGGPT
7 SNSVVKRGDPLYLQPHTVLITKSGVQLIMKSKLPYAIQLQEWLLEEVIPQVLCTGKYDPAIKQREESK
8 QLVTKLIATFTEHTNALQAVVAQKTEELVKKQEFIERIVAIDKQIEAKDLQVTRVMTDLNRMYTGFQET
9 MQKKDEIMQKKDAQVTDLVAKVVDLSDRAVQYPADKRKHPVLCVTRDGTTFITAITGQKTYVENQKHKRNI
10 NVANIVVENIRPNPTVDWNNATDRLQAKRSKRSIVLVRWKKRNNLKIG

```

Figure 4: AcMNPV Animo Acid File Content Example

The first stage of the bacteria class is obtaining the data in the respective *.faa file. In this process there are three parts. The first part is whether the end of the file has been met. If so, the search to get the characters for the bacteria class ends. If not, then the second part comes into effect. If the character identified is a '>', then that line is skipped. This line is skipped because the line does not contain the data required, rather just the name and reference of the amino acid. This can be seen in lines one and five in Figure 4 above. When the line is skipped the characters are read initialise buffer is called. The last part is whether a '\n', or space, character is identified. If not, the process continues to get a character. If a '\n' is identified, then continue buffer is called. This section of the bacteria class can be seen below in Figure 5

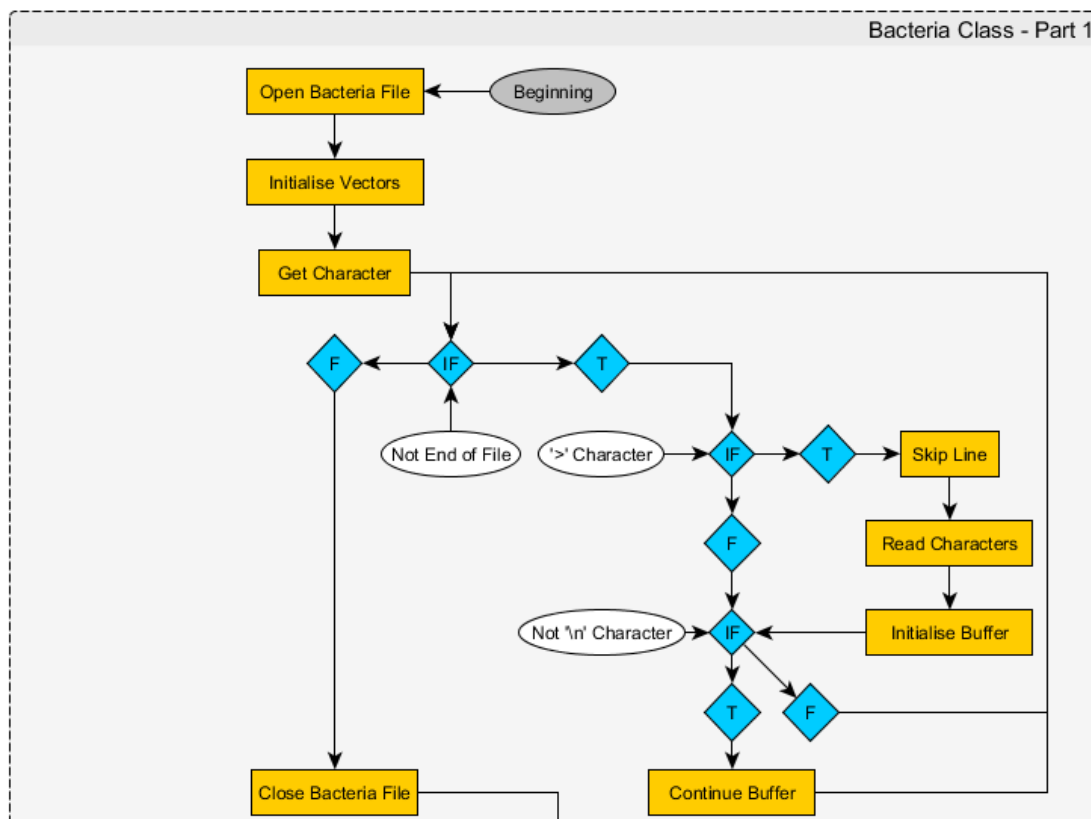


Figure 5: Get Data from *.faa File

When the end of the file has been confirmed the respective file is closed. From here the calculations are performed. Figure 6, shown below, shows the software architecture for this stage up to the point where stochastic is calculated.

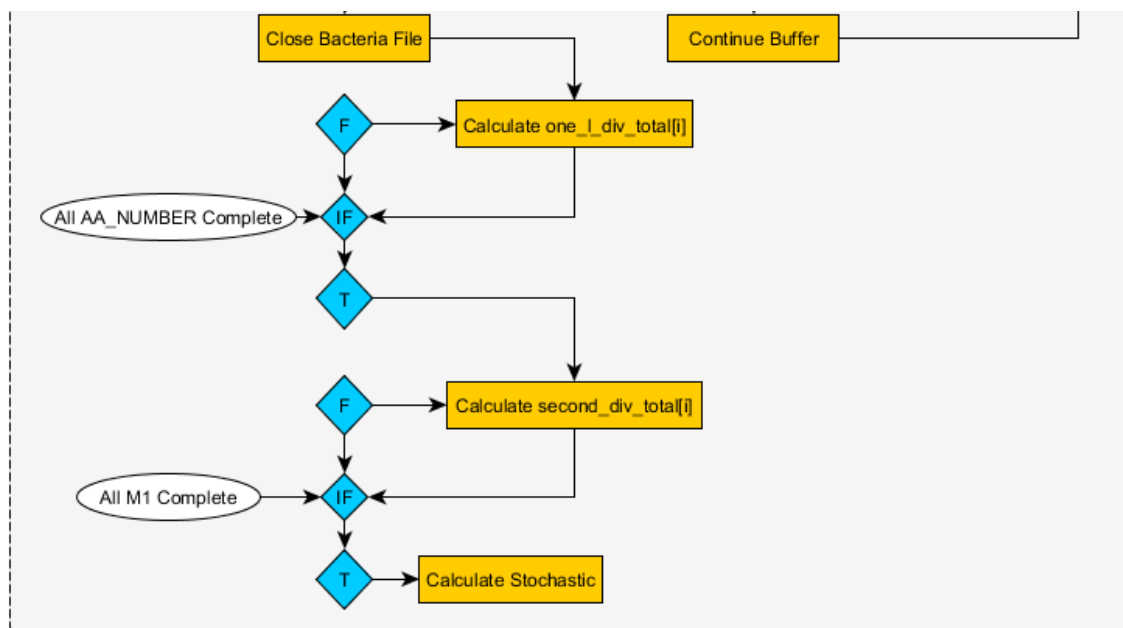


Figure 6: Calculate Variables in Bacteria Class – Part 1

The bacteria class proceeds to calculate other variables which can be seen below in Figure 7.

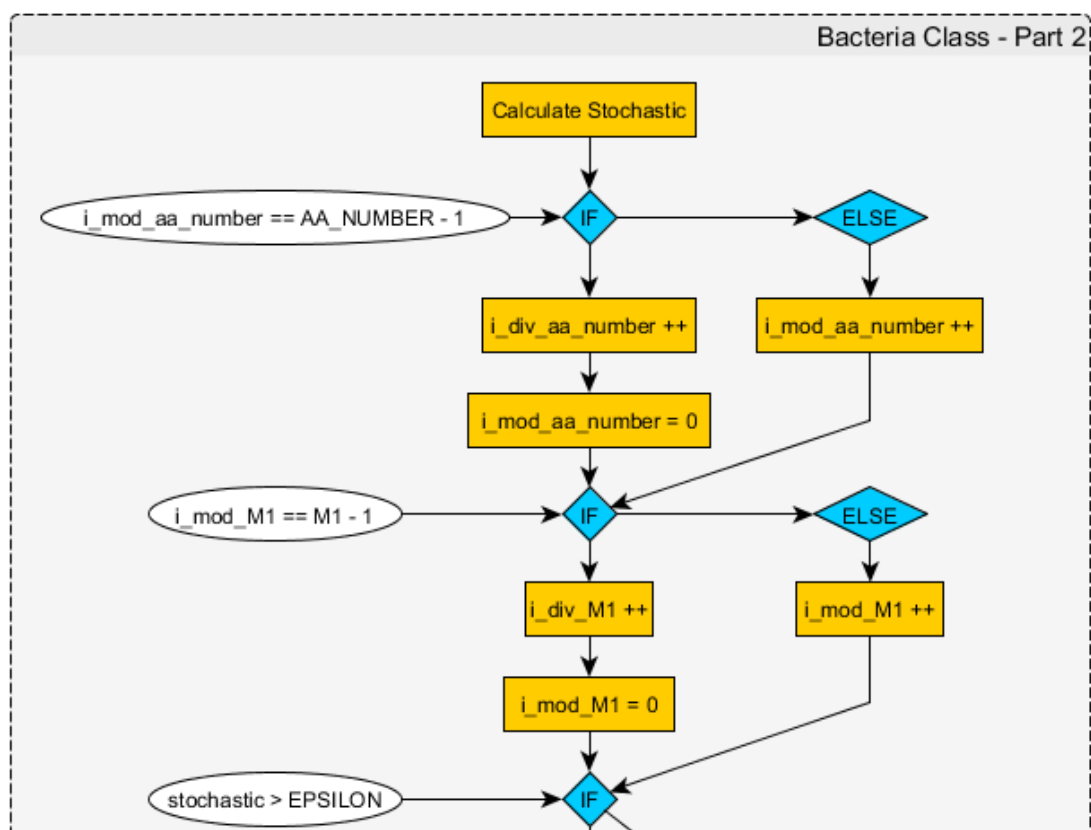


Figure 7: Other Values Calculated in Bacteria Class

The main importance of the bacteria class is the calculated values for $tv[]$ and $ti[]$. These variables are used when comparing two different bacteria. Figure 8 below shows the software architecture for the rest of the bacteria class. The full software architecture diagram for the bacteria class can be found in the appendices. Bacteria Class – Part 1 software architecture can be found in section 6.4 of this report. Bacteria Class – Part 2 software architecture can be found in section 6.5 of this report.

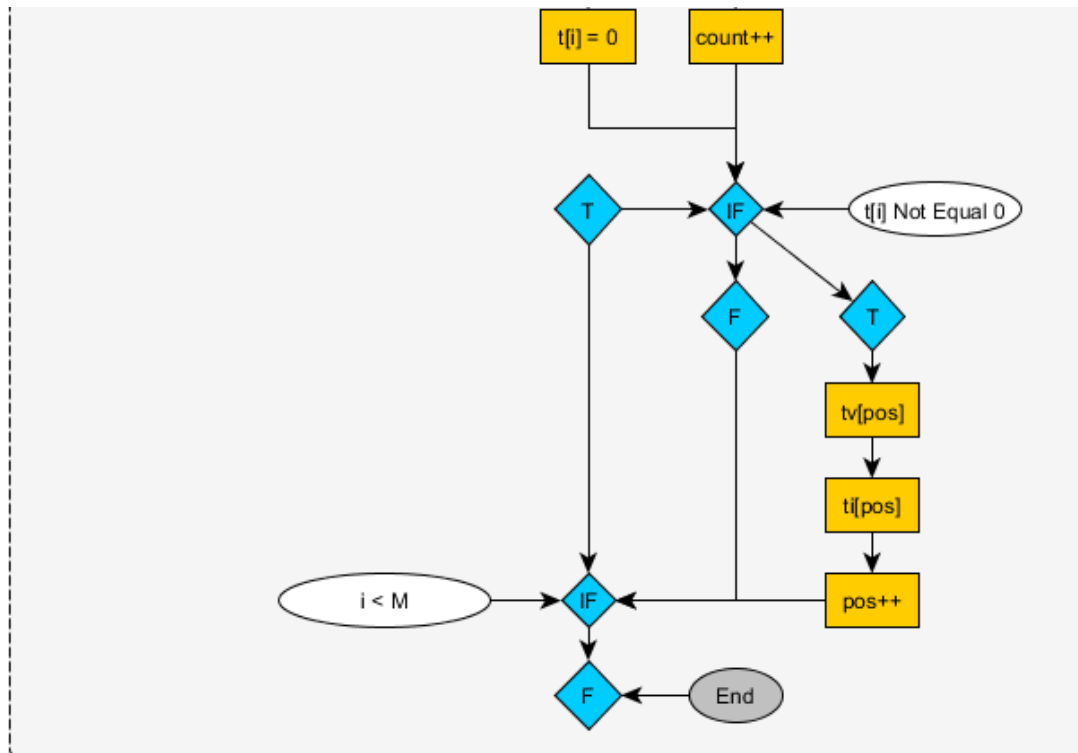


Figure 8: Bacteria Class Ending

1.1.6 Compare Bacteria

The compare bacteria function takes two bacteria as inputs and proceeds to calculate the correlation between the two bacteria based on their vectors, ti and tv variables. In this application, the two bacteria that are taken are two different bacteria. This occurs in the CompareAllBacteria function which includes a nested loop to ensure that every possible unique pair of bacteria are compared. The detail software architecture for the CompareBacteria function can be seen in Figure 9 on the next page and in section 6.6 of the appendices in this report.

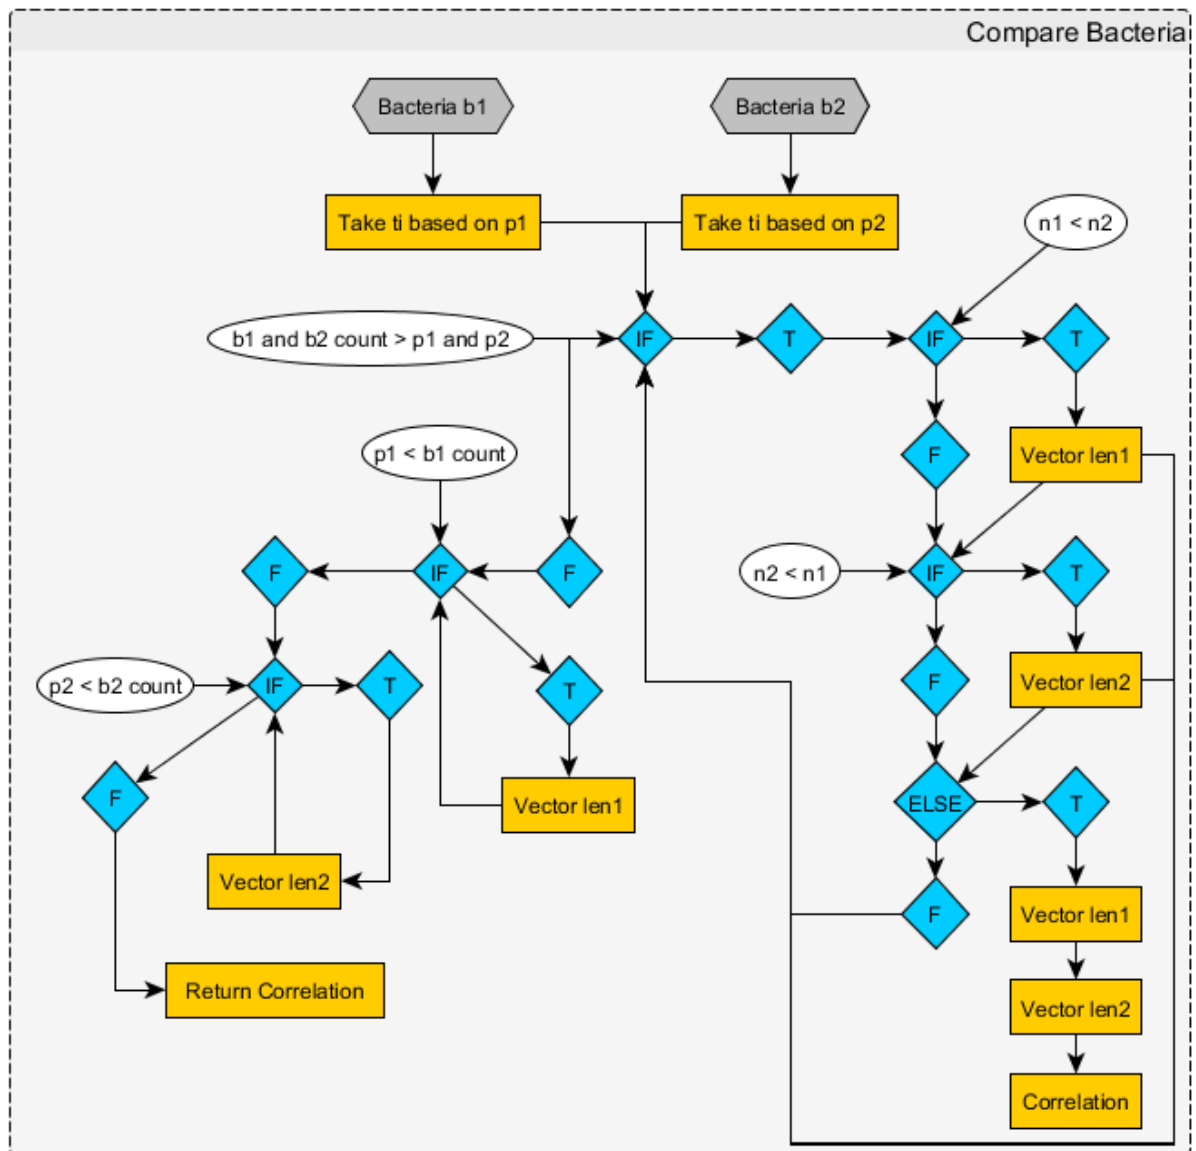


Figure 9: Detail View of Compare Bacteria

2 Bioinformatics – Genome Similarity Using Frequency Vectors

This section of the report discusses what parts of the original code that were analysed. This includes sections of code that resulted in success and failure to be parallelised and what parts of the application can be scalable. This section also discusses the software used and techniques during this analysis.

2.1 Software, Compiler, Hardware and Techniques

2.1.1 Software

The integrated development environment (IDE) used when analysing the application and its code was Microsoft's Visual Studio 2017. This IDE's diagnostic tools were used to identify which section of code took the most amount of time to process. Furthermore, the diagnostic tools show the memory and CPU usage that is being used by the application in real-time. This assisted in analysing for-loop parallelisation as to how much memory was being used and the difference between the original and parallelised version of code's CPU usage.

OpenMP 2.0 is included as part of Visual Studio 2017. This was enabled and chosen for the following reasons. Its simplicity, ease of use and integration in the IDE makes it a wonderful choice when parallelising for-loops and multi-threading other components of the application. The second reason is that I have not used the OpenMP API before. The third reason, OpenMP performed faster results than using POSIX threads. Unfortunately, OpenMP does not perform while-loop parallelisation. While-loops can be parallelised using OpenMP tasks, but this was found to provide no further performance gain and all while loops did not benefit from this.

2.1.2 Compiler

The compiler used was the default compiler that Microsoft has included in Visual Studio 2017. This version of the compiler was Visual C++ 14.1.

2.1.3 Hardware

In this report two machines were used. Both machines use the Windows 10 OS and are running on High Performance mode.

The machine used, named miniTOP, has an i7-4710MQ Intel processor operating at 2.50GHz to 3.50GHz with 16.0 GB of DDR3 RAM. This CPU has four physical cores and eight threads. This machine does not overclock the CPU. This machine's RAM clock rate is at 1600 MHz.

2.1.4 Techniques

OpenMP's sections, parallel and for techniques were used to make the program complete within a faster time frame during testing. Sections to perform two different tasks at the same time or splitting sections of work. Parallel and for were used to parallelise for-loops. Other techniques attempted were removing a nested for-loop for a while loop. Utilising an array to store the correlation variables before printing to the console in order.

2.2 Analysis of Code and Parallelisation Process

2.2.1 Successful Parallelisation

Most of the processing time and memory being used in this application is when the CompareAllBacteria function is called from main. This function contains one for-loop and a nested for-loop. These for-loops call the bacteria class to be filled with *.faa file data and compare bacteria calculates the correlation of the two bacteria. This can be seen in Figure 10 below.

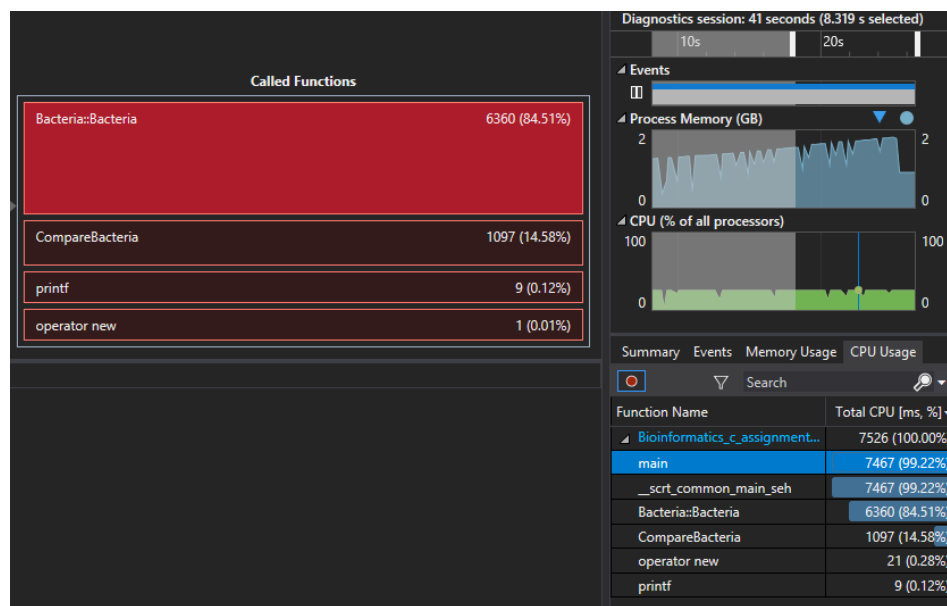


Figure 10: Total CPU (ms, %) of Functions in Bioinformatics – Genome Similarity Using Frequency Vectors Application

The largest components are the bacteria class when files are being loaded and when two bacteria are being compared.

The first for-loop loads each *.faa file into its respective bacteria class. The original code loads each of these files sequentially, one at a time, when there are forty-one files to be loaded. This also fills data into each bacteria class sequentially as well for the respective bacteria file loaded. This for-loop could be parallelised to load multiple bacteria files on numerous threads. This section of code can be found in section 6.1 of the appendix, lines 239 to 243.

The nested for-loop inputs which two bacteria should be compared. This then calculates the correlation between two different bacteria. This nested loop then prints the result of the two bacteria to the console. This nested-loop sequentially calculates and prints each result to the console. This nested-for loop could be parallelised by calculating the correlation of two different bacteria over numerous threads and having them print to the console. However, this could be printed out of order which would lead to a disorganised console output. This section of code can be found in section 6.1 of the appendix, lines 245 to 252.

A solution to keep order but keeping parallelisation would be to load the correlated results into an array and then use the array to print the correlated results in sequential order. This will require changing of the code.

These for-loops should be scalable based on the number of bacteria to be loaded and compared. This will be shown further in section 4 of this report.

2.2.2 Further Analysed Parallelisation Attempts

While most of the run-time is performed in the CompareAllBacteria function, there are other for-loops, while-loops and other sections of code that were analysed to see if there would be any performance increase. However, this was not the case for the following.

2.2.2.1 Two Functions at the Same Time in Main

In the main function of the application, there are two functions that are called. Both functions were found to have a low, if negligible, computational time. However, the ReadInputFile function does not have any variables that are dependant of the previous called Init function. It would be interesting to see if there is a time reduction when running both functions at the same time on two different threads instead of the sequential process of Init being called then ReadInputData. However, this did not provide any faster computation. This is because the code in Init and ReadInputFile proved to have insignificant computation time.

2.2.2.2 Parallelising Bacteria Class

In the bacteria class there are four for-loops. The first for-loop, at lines 108 to 109 at section 6.1 of this report, parallelisation resulted with no additional performance gain. The second for-loop, at lines 112 to 113 at section 6.1 of this report, parallelisation resulted with no additional performance gain.

The third for-loop, at lines 117 to 118 at section 6.1 of this report, OpenMP's parallel and for techniques resulted in different outcomes. The parallel for (combined) parallelisation technique resulted in a memory access violation for the variable p2 calculation. The for parallelisation technique resulted in an infinite loop where nothing is being processed according to Visual Studio's diagnostic tools. The parallel parallelisation technique resulted the application printing the correct results to the console, but at no reduce time.

The fourth for-loop, at lines 158 to 166 at section 6.1 of this report, also had different outcomes when using the OpenMP for-loop parallelisation techniques. Parallel and for combined resulted with incorrect correlation values being printed to the console. The for technique caused a never-ending loop where the processor did no computational work. The parallel technique resulted with a working application printing the correct results, but this did not lower the time spent.

After further analysis with the calculations and variable dependency in the bacteria class, OpenMP sections were utilised to split up the work load over numerous threads. However, this continued to cause a memory access violation at locations at the second for-loop, lines 112 to 113 section 6.1 of this report. This resulted in accepting that the bacteria class could not have its work load split into pieces to reduce to time taken for the application.

2.2.2.3 Parallelising Comparing Two Bacteria

In the CompareBacteria function there is a while-loop at lines 195 to 219 in section 6.1 in the appendices of this report. This while-loop was tested with OpenMP's parallel technique and resulted in a worse performance time for the application. OpenMP's parallel technique was also used on the while-loops from lines 220 to 231 in section 6.1 of this report. The result had no changes to the time taken for the application. These while loops were then split into OpenMP sections as they do not have any related variable or data dependencies with the exception to the variable count. This resulted in the application to crash while running the Debug model. No warnings displayed, or exceptions caught. In the Release model, an infinite loop occurs.

3 Code Implemented

This section of the report includes the code from the original Bioinformatics – Genome Similarity Using Frequency Vectors application in comparison to the final parallelised version. These code changes will start from the top most changes down to the bottom of the code.

3.1 Code Changes in Compare All Bacteria Function

The original version of code for the CompareAllBacteria function can be seen below in Figure 11. Below Figure 11, Figure 12 can be seen which includes the code in the high-performance version.

```

236 void CompareAllBacteria()
237 {
238     Bacteria** b = new Bacteria*[number_bacteria];
239     for (int i = 0; i < number_bacteria; i++)
240     {
241         printf("load %d of %d\n", i + 1, number_bacteria);
242         b[i] = new Bacteria(bacteria_name[i]);
243     }
244
245     for (int i = 0; i < number_bacteria - 1; i++)
246         for (int j = i + 1; j < number_bacteria; j++)
247         {
248             printf("%2d %2d -> ", i, j);
249             double correlation = CompareBacteria(b[i], b[j]);
250             printf("%.20lf\n", correlation);
251         }
252 }

```

Figure 11: Original Code for Compare All Bacteria Function

```

280 void CompareAllBacteria() {
281     Bacteria** b = new Bacteria*[number_bacteria];
282     /*
283     [8] Success with major improvement
284     */
285     #pragma omp parallel for
286     for (int i = 0; i < number_bacteria; i++) {
287         printf("load %d of %d\n", i + 1, number_bacteria);
288         b[i] = new Bacteria(bacteria_name[i]);
289     }
290
291     /*
292     [9]
293     */
294     #pragma omp parallel for
295     for (int i = 0; i < number_bacteria - 1; i++) {
296         /*
297         [10]
298         */
299         for (int j = i + 1; j < number_bacteria; j++) {
300             double correlation = CompareBacteria(b[i], b[j]);
301             printf("%2d %2d -> %.20lf\n", i, j, correlation);
302         }
303     }
304 }

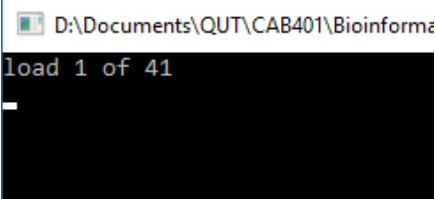
```

Figure 12: Parallelised Code for Compare All Bacteria Function

The complete version of the code for the original version of the application can be found in section 6.1 of the appendices. The complete version of the code for the parallelised version of the application can be found in section 6.2 of the appendices. In Figure 12, there are missing lines of code. These lines were removed as they are commented notes taken during the analysis and parallelisation process.

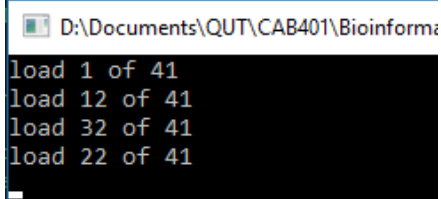
The lines of code in Figure 12, lines 289 and 337, are the OpenMP parallel and for functionality that OpenMP can perform. OpenMP parallel spawns a group of threads. OpenMP for divides the loop iterations between the spawned threads. Both the 'parallel' and 'for' techniques offered by OpenMP were combined. Through testing they provided the quickest time to achieving that all the *.faa files were loaded, all respective data was used to create each respective bacteria class, all bacteria comparisons were calculated, and those results are printed to the console.

Below in Figures 13 and 14 it can be seen what happens in the first second of the two versions of the application when it starts. In Figure 13, the first file is being loaded. In Figure 14, four bacteria *.faa files are being read and those classes are being created for bacteria one, twelve, thirty-two and twenty-two. This is because the test was running on a four core, four read capable CPU so only a maximum of four files can be read at any one time when using this machine. However, on a four core, eight thread capable CPU, a maximum of eight files can be read. This can be seen in Figure 15 on the next page.



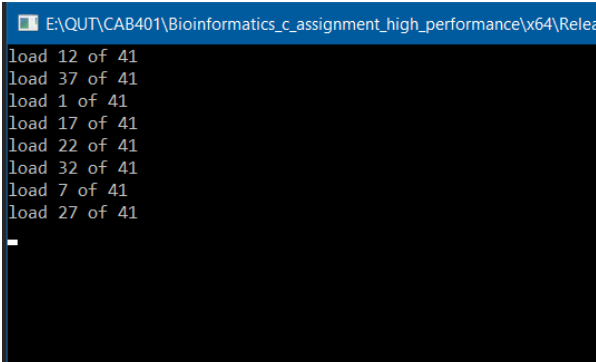
```
D:\Documents\QUT\CAB401\Bioinforma:
load 1 of 41
```

Figure 13: Original Loading Bacteria



```
D:\Documents\QUT\CAB401\Bioinforma:
load 1 of 41
load 12 of 41
load 32 of 41
load 22 of 41
```

Figure 14: Parallel Loading Bacteria Four Threads Max



```
E:\QUT\CAB401\Bioinformatics_c_assignment_high_performance\x64\Relea
load 12 of 41
load 37 of 41
load 1 of 41
load 17 of 41
load 22 of 41
load 32 of 41
load 7 of 41
load 27 of 41
```

Figure 15: Parallel Loading Bacteria Eight Threads Max

Using OpenMP's parallel and for techniques for the for-loop at lines 289 to 293, shown at Figure 12, multiple files can be read, and bacteria classes created at the same time. The obvious limitation here is the read speed of the disc that the files are stored on. However, the use of OpenMP allows for the application to be scalable. This scalability is dependent of the

maximum number of threads of a CPU, the clock rate of the CPU, enough memory and the number of bacteria being compared. Below, in Figures 16 and 17, the amount of memory being used increases with each bacteria file being read and class created.

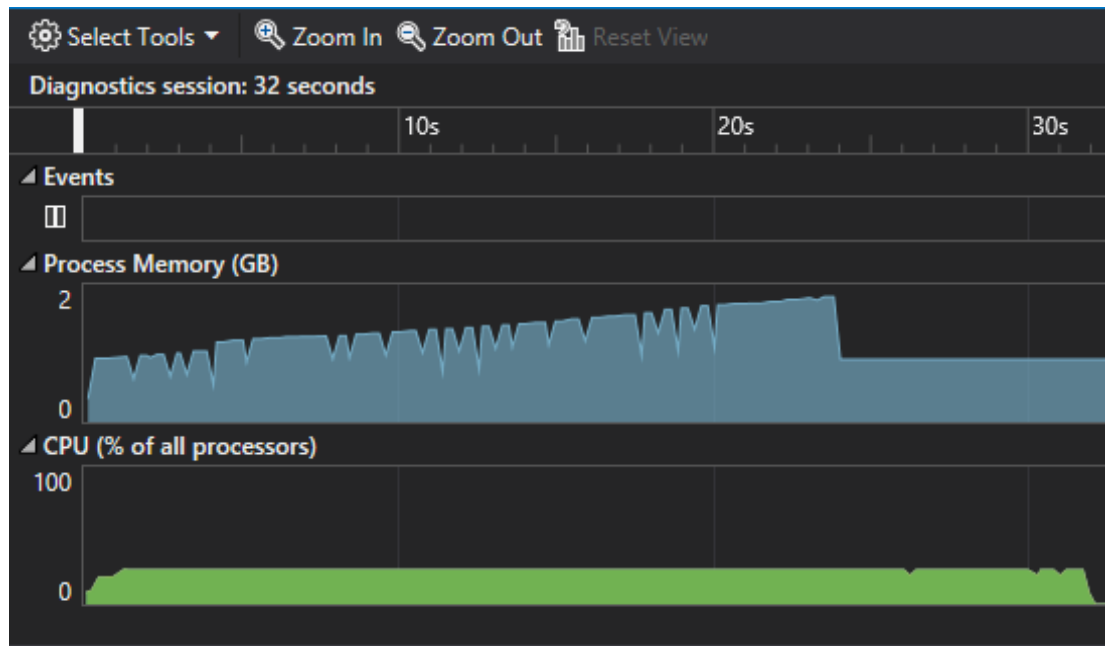


Figure 16: Sequential Application Memory and CPU Usage

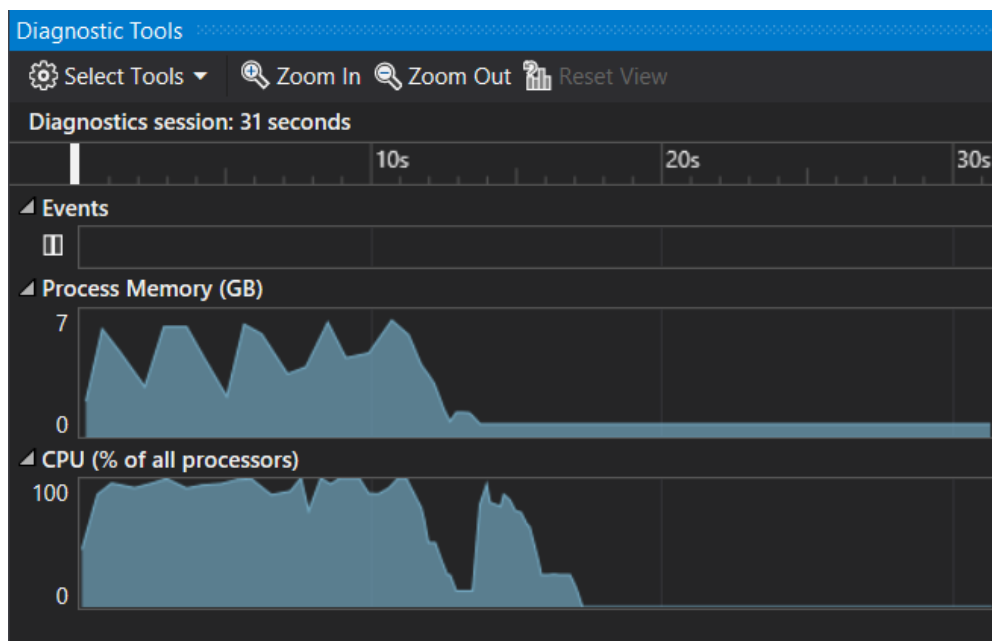


Figure 17: Parallelised Application Memory and CPU Usage

For the parallelisation of the nested for-loop, from lines 337 to 346 in Figure 12, the correlation printed will be in a mixed order. While, the results are still the same, only the order of the printed results are different. If the user wished to have the correlation printed in numerical order, bacteria 0 and bacteria 1 to bacteria 39 and bacteria 40, but still wish for a faster performance time using OpenMP then Figure 18 on the next page can be used instead.

```

295     /*
296     [9]
336     */
337     #pragma omp for
338     for (int i = 0; i < number_bacteria - 1; i++) {
339         /*
340         [10]
341         */
342         for (int j = i + 1; j < number_bacteria; j++) {
343             double correlation = CompareBacteria(b[i], b[j]);
344             printf("%2d %2d -> %.201f\n", i, j, correlation);
345         }
346     }
347 }

```

Figure 18: Parallelised Nested For-Loop in Order

The difference between Figure 12 and Figure 18 is that the OpenMP statement removed the parallel functionality. This added a few seconds of extra time computational time to the application, but the printed correlation and bacteria pairs are printed to the console in order. This order difference can be seen in Figures 19 and 20 below.

```

3 32 -> 0.00173857472560830185
3 33 -> 0.00119861580795493547
3 34 -> 0.00132066757407171918
3 35 -> 0.00115326991777306800
3 36 -> 0.00126922941777803675
3 37 -> 0.00155134347347462751

```

Figure 19: Correlation in Order

```

3 32 -> 0.00173857472560830185
28 30 -> 0.00208711780367286511
3 33 -> 0.00119861580795493547
28 31 -> 0.00126214769015896462
17 20 -> 0.00174925382121538398
28 32 -> 0.00192751640765809564

```

Figure 20: Correlation not in Order

This version of the code was not taken though as the purpose of this assignment was to take a sequential program that could be parallelised and then parallelised to create the same results but in a shorter time frame that still produces the same results.

4 Results

In this section of the report, the correlation and times taken to run the application, both the original and high-performance models, on two different machines are discussed.

4.1 Time Taken for Execution of Sequential and Parallelised Applications

The systems specification for this machine can be found in section 2.1.3 of this report.

Table 1 below shows the time to complete the original and the parallelised versions of the Bioinformatics – Genome similarity using Frequency Vectors application. The application was executed five times and the average mean time was calculated. The parallelised version had each OpenMP line of code set to the specified number of threads below.

Application Version	Time Taken (Seconds)					
	1st	2nd	3rd	4th	5th	Average
Original	28 secs	29 secs	28 secs	28 secs	28 secs	28.2 secs
High Performance (1 Thread)	29 secs	29 secs	29 secs	29 secs	29 secs	29.0 secs
High Performance (2 Threads)	20 secs	20 secs	20 secs	20 secs	20 secs	20.0 secs
High Performance (3 Threads)	18 secs	17 secs	17 secs	17 secs	17 secs	16.2 secs
High Performance (4 Threads)	16 secs	17 secs	17 secs	17 secs	17 secs	16.8 secs
High Performance (5 Threads)	16 secs	16 secs	16 secs	16 secs	16 secs	16.0 secs
High Performance (6 Threads)	16 secs	16 secs	16 secs	16 secs	16 secs	16.0 secs
High Performance (7 Threads)	16 secs	16 secs	16 secs	16 secs	16 secs	16.0 secs
High Performance (8 Threads)	16 secs	16 secs	16 secs	16 secs	16 secs	16.0 secs

Table 1: Sequential and Different Thread Count for Parallelised Application Run-Time

From the results shown above I cannot state whether the application is limited to four threads or if the test results were limited due to the number of physical cores available on the hardware used. A machine with more than four physical cores is required to prove the previous statement.

4.2 Profiling Reports

Figure 21 below is a CPU profiling report created in Visual Studio 2017 showing the CPU usage of the sequential application and percentage of work for functions during run-time.

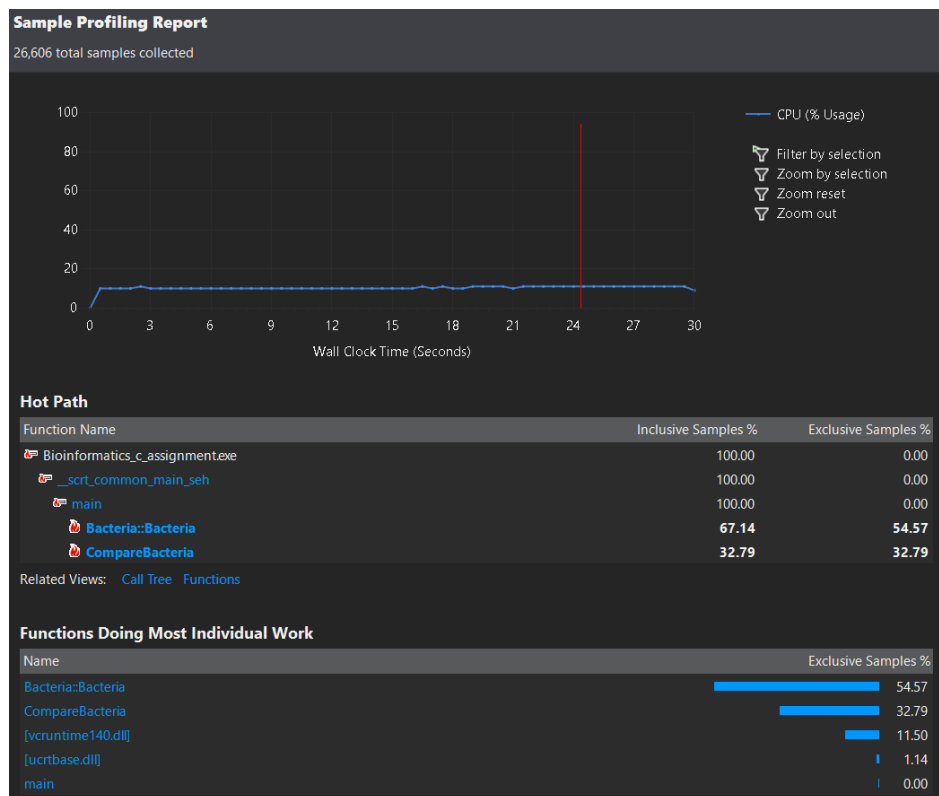


Figure 21: Sequential CPU Profiler

Figure 22 below is a CPU profiling report created in Visual Studio 2017 showing the CPU usage of the parallelised application and percentage of work for functions during run-time.

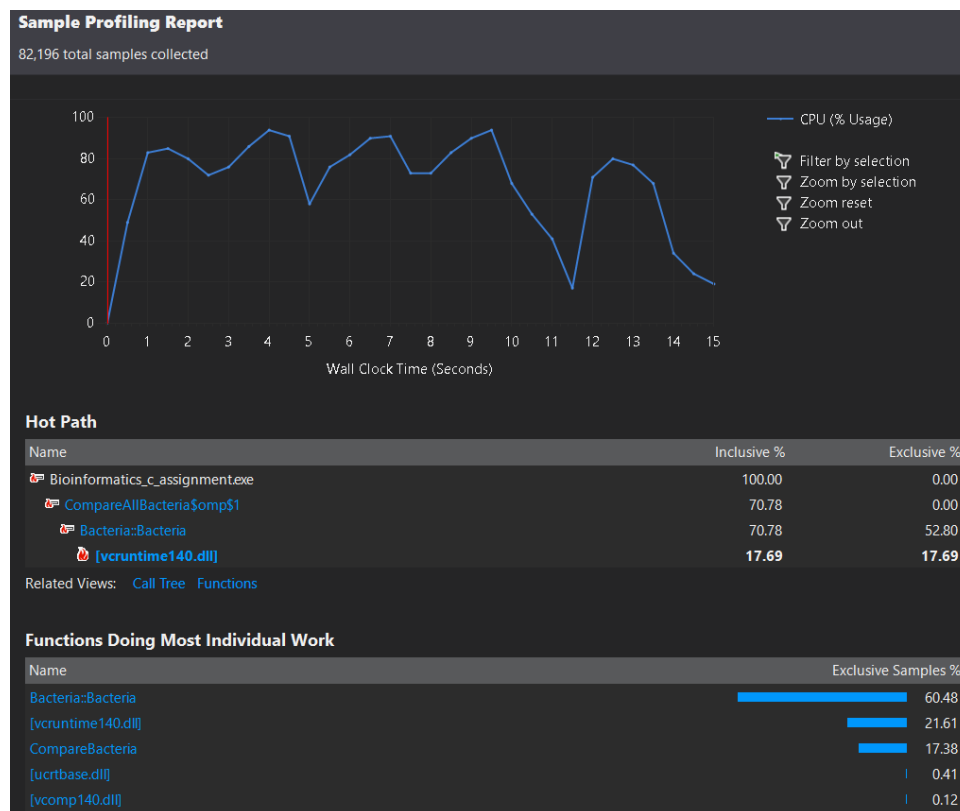


Figure 22: Parallelised CPU Profiler

4.3 Speed Up Graph

The graph below shows the speed up of the parallelised version of the application compared to the sequential application based on the time difference and the number of threads being used.

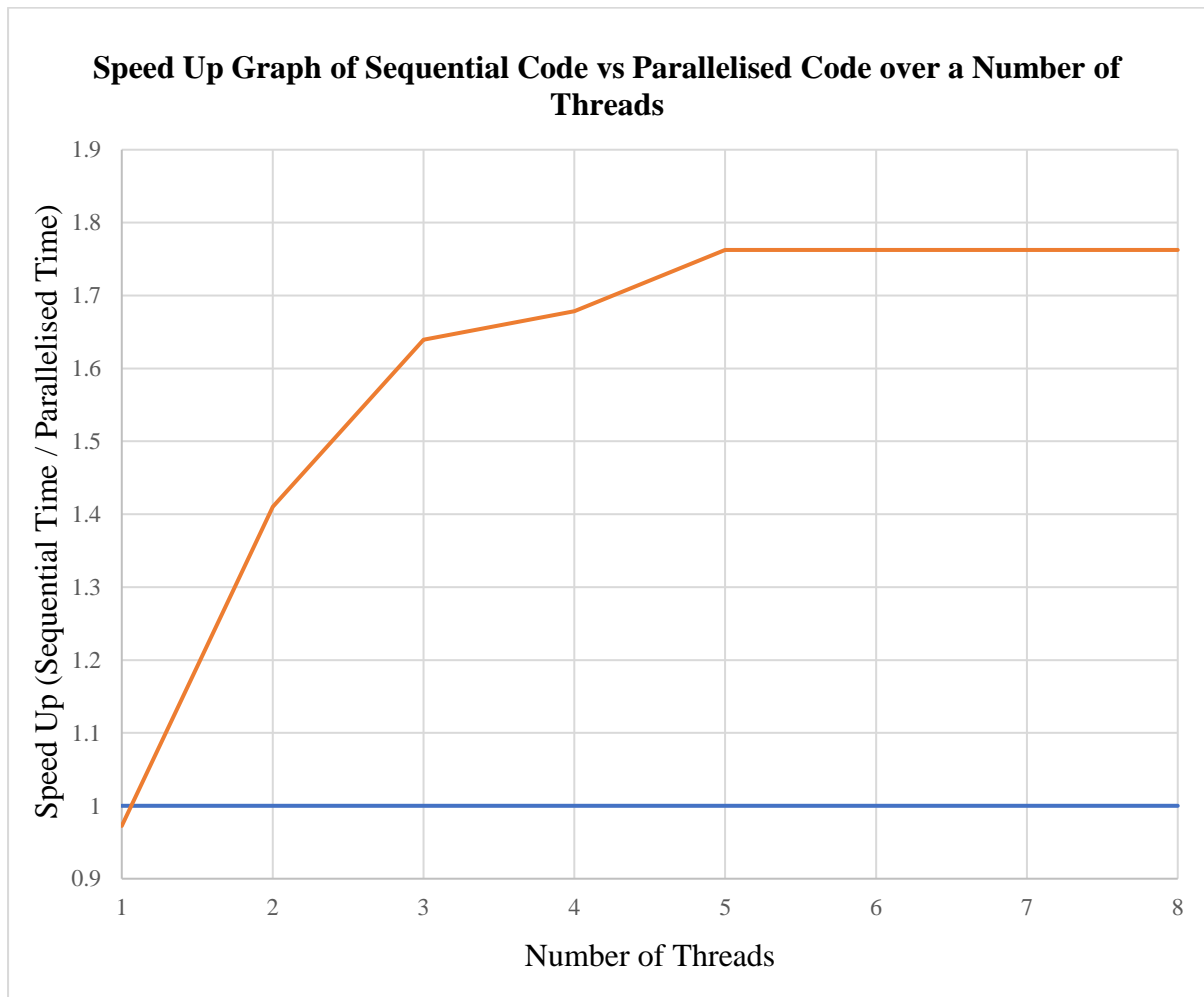


Figure 23: Speed Up Graph of Sequential vs Parallelised Over Threads

5 Reflection

Originally the Bioinformatics – Genome similarity using Frequency Vectors application was supposed to be parallelised using MPI, mpich3 or openMP, on my Raspberry Pi Beowulf cluster. This was scrapped because the application immediately gets a segmentation fault. Instead of fiddling around to get the code for the application to run properly I opted-out for running the application on through Visual Studio on a Windows OS machine so I could get the application parallelised. In terms of what I have learned, I usually always utilised Code::Blocks and the mingw 32 and 64 bit compilers when I program in C/C++. This is mainly because the IDE and compiler is free and widely used. However, during this semester, the CAB401 and IGB381 courses showed me how useful Visual Studio can be. The profiler and diagnostic tools allows me to see what functions and code could be further improved. Much easier than slowly debugging your way through the application manually. Another thing I have learnt is how simple OpenMP is to parallelise code with. I've always used POSIX threads when I wanted to parallelise for-loops or create multiple simultaneous executions of work. While there is a difference between the two, low-level vs high-level portable multiprocessing paradigms, it was a good experience to utilise and learn what OpenMP has to offer and how it works.

I believe my attempt at analysing the original code and going beyond trying different techniques to get the application to perform faster was almost a success. The one component that bothers me is the nested for-loop in the CompareAllBacteria function. An issue was encountered here that I could not fix.

5.1 Issues Encountered

I attempted to store the correlated data into a vector, dynamic array by splitting up the objectives. These objective were to calculate the correlation and printing to the console. This code can be seen in Figure 24 below.

```
std::vector<double> correlation;
int vectorSize = (number_bacteria * number_bacteria) - number_bacteria;
correlation.resize(vectorSize);

#pragma omp parallel for
for (int i = 0; i < number_bacteria - 1; i++) {
    for (int j = i + 1; j < number_bacteria; j++) {
        correlation[(i * 10) + j] = CompareBacteria(b[i], b[j]);
    }
}

for (int i = 0; i < number_bacteria - 1; i++) {
    for (int j = i + 1; j < number_bacteria; j++) {
        printf("%2d %2d -> %.20lf\n", i, j, correlation[(i * 10) +
j]);
    }
}
```

Figure 24: Restructured Nested For-Loop Attempt

The result of this code sometimes provided the correct correlation values. At other times I found the correct correlation value, being returned by the function CompareBacteria, was

changed when it was stored in the vector. This produced incorrect results. The purpose of this was to allow for the exact run-time to be achieved but with the correlated values printed in order. Splitting these objectives into two components allowed for the CompareBacteria to be parallelised providing a faster run-time. Then a sequential for-loop to print out the results in order.

It annoys me to know that I could have improved the ease of use, printing correlation in order, especially when a colleague of mine in CAB401 did the same thing and it worked for him.

As discussed in section 2.2.2, of this report, it was found that many data dependencies in other functions prevented further changes and parallelism for this application. At times the incorrect correlation between bacteria would occur. Further problems such as infinite loops and application crashes were encountered to further improve the application. However, solutions to these problems provided parallelisation that had no effect on the applications run-time making the parallelisation of multiple sections of code running at the same time on different threads pointless.

For more notes during testing the high-performance version of the Bioinformatics – Genome similarity using Frequency Vectors application can be found commented in the *.cpp file of the Visual Studio project.

6 Appendix

6.1 Original Code

The original code for the Bioinformatics – Genome similarity using Frequency Vectors (C++) application was provided by Dr Wayne Kelly on the CAB401 High Performance and Parallel Computing QUT Blackboard page. A copy of the original code can be seen below. All bacteria files and the Visual Studio 2017 project can be found at the link here https://github.com/Starwolf-001/CAB401/tree/master/Bioinformatics_c_assignment.

```

1      #include <stdio.h>
2      #include <string.h>
3      #include <time.h>
4      #include <math.h>
5      #include <iostream>
6
7      int number_bacteria;
8      char** bacteria_name;
9      long M, M1, M2;
10     short code[27] = { 0, 2, 1, 2, 3, 4, 5, 6, 7, -1, 8, 9, 10, 11, -1, 12,
11                      13, 14, 15, 16, 1, 17, 18, 5, 19, 3 };
12     #define encode(ch)      code[ch-'A']
13     #define LEN            6
14     #define AA_NUMBER      20
15     #define EPSILON        1e-010
16
17     void Init()
18     {
19         M2 = 1;
20         for (int i = 0; i<LEN - 2; i++) // M2 = AA_NUMBER ^ (LEN-2);
21             M2 *= AA_NUMBER;
22         M1 = M2 * AA_NUMBER;           // M1 = AA_NUMBER ^ (LEN-1);
23         M = M1 * AA_NUMBER;           // M = AA_NUMBER ^ (LEN);
24     }
25
26     class Bacteria
27     {
28     private:
29         long* vector;
30         long* second;
31         long one_l[AA_NUMBER];
32         long indexs;
33         long total;
34         long total_l;
35         long complement;
36
37         void InitVectors()
38         {
39             vector = new long[M];
40             second = new long[M1];
41             memset(vector, 0, M * sizeof(long));
42             memset(second, 0, M1 * sizeof(long));
43             memset(one_l, 0, AA_NUMBER * sizeof(long));
44             total = 0;
45             total_l = 0;
46             complement = 0;
47         }
48
49         void init_buffer(char* buffer)
50         {
51             complement++;
52             indexs = 0;

```



```
53         for (int i = 0; i<LEN - 1; i++)
54         {
55             short enc = encode(buffer[i]);
56             one_l[enc]++;
57             total_l++;
58             indexs = indexs * AA_NUMBER + enc;
59         }
60         second[indexs]++;
61     }
62
63     void cont_buffer(char ch)
64     {
65         short enc = encode(ch);
66         one_l[enc]++;
67         total_l++;
68         long index = indexs * AA_NUMBER + enc;
69         vector[index]++;
70         total++;
71         indexs = (indexs % M2) * AA_NUMBER + enc;
72         second[indexs]++;
73     }
74
75 public:
76     long count;
77     double* tv;
78     long *ti;
79
80     Bacteria(char* filename)
81     {
82         FILE * bacteria_file = fopen(filename, "r");
83         InitVectors();
84
85         char ch;
86         while ((ch = fgetc(bacteria_file)) != EOF)
87         {
88             if (ch == '>')
89             {
90                 // skip rest of line
91                 while (fgetc(bacteria_file) != '\n');
92                 char buffer[LEN - 1];
93                 fread(buffer, sizeof(char), LEN - 1,
94                     bacteria_file);
95                 init_buffer(buffer);
96             }
97             else if (ch != '\n')
98                 cont_buffer(ch);
99         }
100
101         long total_plus_complement = total + complement;
102         double total_div_2 = total * 0.5;
103         int i_mod_aa_number = 0;
104         int i_div_aa_number = 0;
105         long i_mod_M1 = 0;
106         long i_div_M1 = 0;
107
108         double one_l_div_total[AA_NUMBER];
109         for (int i = 0; i<AA_NUMBER; i++)
110             one_l_div_total[i] = (double)one_l[i] / total_l;
111
112         double* second_div_total = new double[M1];
113         for (int i = 0; i<M1; i++)
114             second_div_total[i] = (double)second[i] /
115                 total_plus_complement;
116
117         count = 0;
```

```
115         double* t = new double[M];
116
117         for (long i = 0; i<M; i++)
118         {
119             double p1 = second_div_total[i_div_aa_number];
120             double p2 = one_l_div_total[i_mod_aa_number];
121             double p3 = second_div_total[i_mod_M1];
122             double p4 = one_l_div_total[i_div_M1];
123             double stochastic = (p1 * p2 + p3 * p4) *
total_div_2;
124
125             if (i_mod_aa_number == AA_NUMBER - 1)
126             {
127                 i_mod_aa_number = 0;
128                 i_div_aa_number++;
129             }
130             else
131                 i_mod_aa_number++;
132
133             if (i_mod_M1 == M1 - 1)
134             {
135                 i_mod_M1 = 0;
136                 i_div_M1++;
137             }
138             else
139                 i_mod_M1++;
140
141             if (stochastic > EPSILON)
142             {
143                 t[i] = (vector[i] - stochastic) / stochastic;
144                 count++;
145             }
146             else
147                 t[i] = 0;
148         }
149
150         delete second_div_total;
151         delete vector;
152         delete second;
153
154         tv = new double[count];
155         ti = new long[count];
156
157         int pos = 0;
158         for (long i = 0; i<M; i++)
159         {
160             if (t[i] != 0)
161             {
162                 tv[pos] = t[i];
163                 ti[pos] = i;
164                 pos++;
165             }
166         }
167         delete t;
168
169         fclose(bacteria_file);
170     }
171 };
172
173 void ReadInputFile(char* input_name)
174 {
175     FILE* input_file = fopen(input_name, "r");
176     fscanf(input_file, "%d", &number_bacteria);
177     bacteria_name = new char[number_bacteria];
```

```
178
179     for (long i = 0; i < number_bacteria; i++)
180     {
181         bacteria_name[i] = new char[20];
182         fscanf(input_file, "%s", bacteria_name[i]);
183         strcat(bacteria_name[i], ".faa");
184     }
185     fclose(input_file);
186 }
187
188 double CompareBacteria(Bacteria* b1, Bacteria* b2)
189 {
190     double correlation = 0;
191     double vector_len1 = 0;
192     double vector_len2 = 0;
193     long p1 = 0;
194     long p2 = 0;
195     while (p1 < b1->count && p2 < b2->count)
196     {
197         long n1 = b1->ti[p1];
198         long n2 = b2->ti[p2];
199         if (n1 < n2)
200         {
201             double t1 = b1->tv[p1];
202             vector_len1 += (t1 * t1);
203             p1++;
204         }
205         else if (n2 < n1)
206         {
207             double t2 = b2->tv[p2];
208             p2++;
209             vector_len2 += (t2 * t2);
210         }
211         else
212         {
213             double t1 = b1->tv[p1++];
214             double t2 = b2->tv[p2++];
215             vector_len1 += (t1 * t1);
216             vector_len2 += (t2 * t2);
217             correlation += t1 * t2;
218         }
219     }
220     while (p1 < b1->count)
221     {
222         long n1 = b1->ti[p1];
223         double t1 = b1->tv[p1++];
224         vector_len1 += (t1 * t1);
225     }
226     while (p2 < b2->count)
227     {
228         long n2 = b2->ti[p2];
229         double t2 = b2->tv[p2++];
230         vector_len2 += (t2 * t2);
231     }
232
233     return correlation / (sqrt(vector_len1) * sqrt(vector_len2));
234 }
235
236 void CompareAllBacteria()
237 {
238     Bacteria** b = new Bacteria*[number_bacteria];
239     for (int i = 0; i < number_bacteria; i++)
240     {
241         printf("load %d of %d\n", i + 1, number_bacteria);
```

```
242         b[i] = new Bacteria(bacteria_name[i]);
243     }
244
245     for (int i = 0; i < number_bacteria - 1; i++)
246         for (int j = i + 1; j < number_bacteria; j++)
247         {
248             printf("%2d %2d -> ", i, j);
249             double correlation = CompareBacteria(b[i], b[j]);
250             printf("%.20lf\n", correlation);
251         }
252     }
253
254     int main(int argc, char * argv[])
255     {
256         time_t t1 = time(NULL);
257
258         Init();
259         ReadInputFile(argv[1]);
260         CompareAllBacteria();
261
262         time_t t2 = time(NULL);
263         printf("time elapsed: %d seconds\n", t2 - t1);
264
265         system("pause");
266
267         return 0;
277     }
```

6.2 High Performance and Parallel Code

This code for the Bioinformatics – Genome similarity using Frequency Vectors (C++) application was analysed and tested to provide the correct results in a short time as per the requirement for the CAB401 High Performance and Parallel Computing Assignment. This code can be seen below. All bacteria files, notes taken during testing, code and the Visual Studio 2017 project can be found at the link here https://github.com/Starwolf-001/CAB401/tree/master/Bioinformatics_c_assignment_high_performance. There are lines missing in this code. These missing lines of code are commented notes based on the testing, and analysis when parallelising the original Bioinformatics – Genome similarity using Frequency Vectors application.

```

1      #include <stdio.h>
2      #include <string.h>
3      #include <time.h>
4      #include <math.h>
5      #include <iostream>
6      #include <omp.h>
7
8      int number_bacteria;
9      char** bacteria_name;
10     long M, M1, M2;
11     short code[27] = {0, 2, 1, 2, 3, 4, 5, 6, 7, -1, 8, 9, 10, 11, -1, 12,
12     13, 14, 15, 16, 1, 17, 18, 5, 19, 3};
13     #define encode(ch)      code[ch-'A']
14     #define LEN             6
15     #define AA_NUMBER      20
16     #define EPSILON        1e-010
17
18     void Init() {
19         M2 = 1;
20         for (int i = 0; i < LEN - 2; i++) {    // M2 = AA_NUMBER ^ (LEN-2);
21             M2 *= AA_NUMBER;
22         }
23         M1 = M2 * AA_NUMBER;                  // M1 = AA_NUMBER ^ (LEN-1);
24         M = M1 * AA_NUMBER;                   // M = AA_NUMBER ^ (LEN);
25     }
26
27     class Bacteria {
28     private:
29         long* vector;
30         long* second;
31         long one_l[AA_NUMBER];
32         long indexs;
33         long total;
34         long total_l;
35         long complement;
36
37         void InitVectors() {
38             vector = new long[M];
39             second = new long[M1];
40             memset(vector, 0, M * sizeof(long));
41             memset(second, 0, M1 * sizeof(long));
42             memset(one_l, 0, AA_NUMBER * sizeof(long));
43             total = 0;
44             total_l = 0;
45             complement = 0;
46         }
47
48         void init_buffer(char* buffer) {

```

```
53         complement++;
54         indexs = 0;
55         for (int i = 0; i < LEN - 1; i++) {
56             short enc = encode(buffer[i]);
57             one_l[enc]++;
58             total_l++;
59             indexs = indexs * AA_NUMBER + enc;
60         }
61         second[indexs]++;
62     }
63
64     void cont_buffer(char ch) {
65         short enc = encode(ch);
66         one_l[enc]++;
67         total_l++;
68         long index = indexs * AA_NUMBER + enc;
69         vector[index]++;
70         total++;
71         indexs = (indexs % M2) * AA_NUMBER + enc;
72         second[indexs]++;
73     }
74
75     public:
76         long count;
77         double* tv;
78         long *ti;
79
80         Bacteria(char* filename) {
81             FILE * bacteria_file = fopen(filename, "r");
82             InitVectors();
83
84             char ch;
85             while ((ch = fgetc(bacteria_file)) != EOF) {
86                 if (ch == '>') {
87                     while (fgetc(bacteria_file) != '\n'); // skip
88                     rest of line
89
90                     char buffer[LEN - 1];
91                     fread(buffer, sizeof(char), LEN - 1,
92                         bacteria_file);
93                     init_buffer(buffer);
94                 }
95                 else if (ch != '\n')
96                     cont_buffer(ch);
97             }
98             fclose(bacteria_file);
99
100             long total_plus_complement = total + complement;
101             double total_div_2 = total * 0.5;
102             int i_mod_aa_number = 0;
103             int i_div_aa_number = 0;
104             long i_mod_M1 = 0;
105             long i_div_M1 = 0;
106
107             double one_l_div_total[AA_NUMBER];
108             double* second_div_total = new double[M1];
109             for (int i = 0; i < AA_NUMBER; i++) {
110                 one_l_div_total[i] = (double)one_l[i] / total_l;
111             }
112             for (int i = 0; i < M1; i++) {
113                 second_div_total[i] = (double)second[i] /
114                     total_plus_complement;
115             }
116
117         }
```

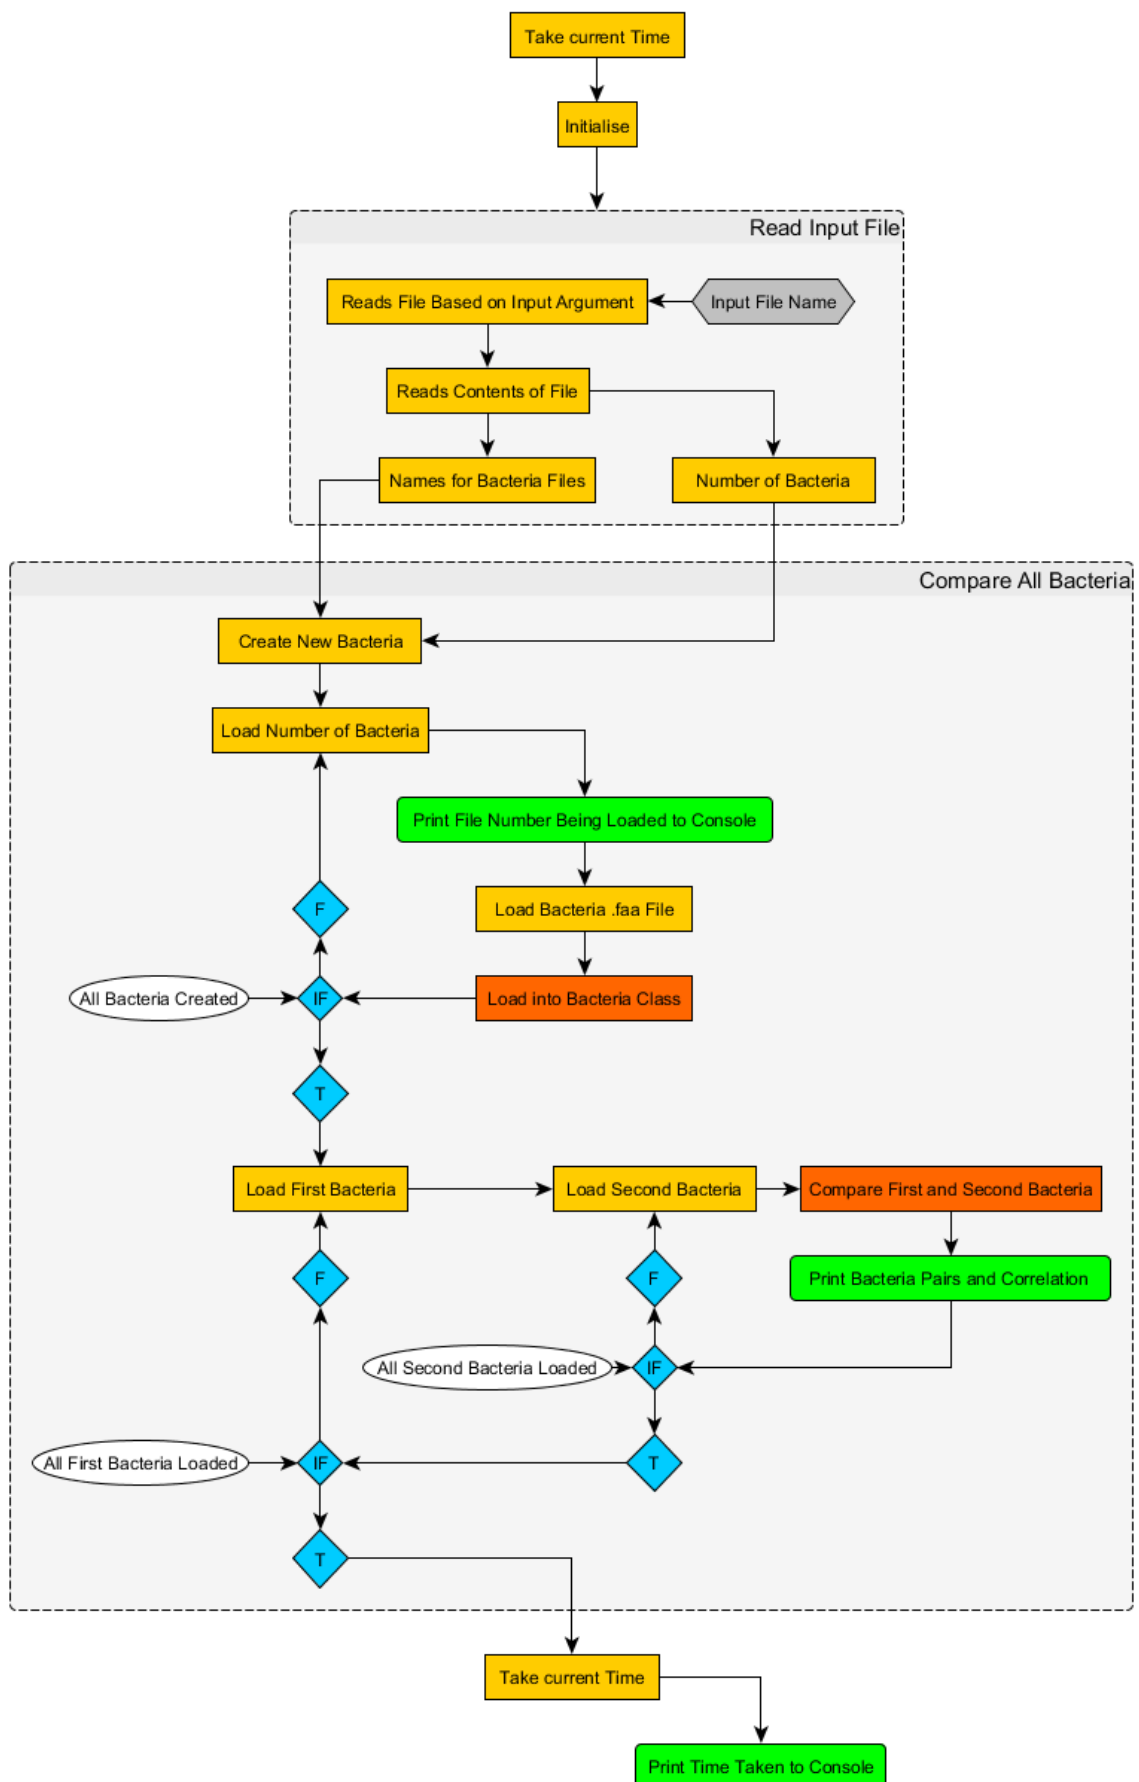
```
142         count = 0;
143         double* t = new double[M];
144
145         for (long i = 0; i < M; i++) {
146             double p1 = second_div_total[i_div_aa_number];
147             double p2 = one_l_div_total[i_mod_aa_number];
148             double p3 = second_div_total[i_mod_M1];
149             double p4 = one_l_div_total[i_div_M1];
150             double stochastic = (p1 * p2 + p3 * p4) *
151                 total_div_2;
152
153             if (i_mod_aa_number == AA_NUMBER - 1) {
154                 i_mod_aa_number = 0;
155                 i_div_aa_number++;
156             }
157             else {
158                 i_mod_aa_number++;
159             }
160
161             if (i_mod_M1 == M1 - 1) {
162                 i_mod_M1 = 0;
163                 i_div_M1++;
164             }
165             else {
166                 i_mod_M1++;
167             }
168
169             if (stochastic > EPSILON) {
170                 t[i] = (vector[i] - stochastic) / stochastic;
171                 count++;
172             }
173             else {
174                 t[i] = 0;
175             }
176         }
177
178         delete second_div_total;
179         delete vector;
180         delete second;
181
182         tv = new double[count];
183         ti = new long[count];
184
185         int pos = 0;
186         for (long i = 0; i < M; i++) {
187             if (t[i] != 0) {
188                 tv[pos] = t[i];
189                 ti[pos] = i;
190                 pos++;
191             }
192         }
193         delete t;
194     }
195 };
196
197 void ReadInputFile(char* input_name) {
198     FILE* input_file = fopen(input_name, "r");
199     fscanf(input_file, "%d", &number_bacteria);
200     bacteria_name = new char*[number_bacteria];
201     for (long i = 0; i < number_bacteria; i++) {
202         bacteria_name[i] = new char[20];
203         fscanf(input_file, "%s", bacteria_name[i]);
204         strcat(bacteria_name[i], ".faa");
205     }
206 }
```

```
225         fclose(input_file);
226     }
227
228     double CompareBacteria(Bacteria* b1, Bacteria* b2) {
229         double correlation = 0;
230         double vector_len1 = 0;
231         double vector_len2 = 0;
232         long p1 = 0;
233         long p2 = 0;
234         while (p1 < b1->count && p2 < b2->count) {
235             long n1 = b1->ti[p1];
236             long n2 = b2->ti[p2];
237             if (n1 < n2) {
238                 double t1 = b1->tv[p1];
239                 vector_len1 += (t1 * t1);
240                 p1++;
241             }
242             else if (n2 < n1) {
243                 double t2 = b2->tv[p2];
244                 p2++;
245                 vector_len2 += (t2 * t2);
246             }
247             else {
248                 double t1 = b1->tv[p1++];
249                 double t2 = b2->tv[p2++];
250                 vector_len1 += (t1 * t1);
251                 vector_len2 += (t2 * t2);
252                 correlation += t1 * t2;
253             }
254         }
255         while (p1 < b1->count) {
256             long n1 = b1->ti[p1];
257             double t1 = b1->tv[p1++];
258             vector_len1 += (t1 * t1);
259         }
260         while (p2 < b2->count) {
261             long n2 = b2->ti[p2];
262             double t2 = b2->tv[p2++];
263             vector_len2 += (t2 * t2);
264         }
265         return correlation / (sqrt(vector_len1) * sqrt(vector_len2));
266     }
267
268     void CompareAllBacteria() {
269         Bacteria** b = new Bacteria*[number_bacteria];
270         #pragma omp parallel for
271         for (int i = 0; i < number_bacteria; i++) {
272             printf("load %d of %d\n", i + 1, number_bacteria);
273             b[i] = new Bacteria(bacteria_name[i]);
274         }
275
276         #pragma omp parallel for
277         for (int i = 0; i < number_bacteria - 1; i++) {
278             for (int j = i + 1; j < number_bacteria; j++) {
279                 double correlation = CompareBacteria(b[i], b[j]);
280                 printf("%2d %2d -> %.20lf\n", i, j, correlation);
281             }
282         }
283     }
284
285     int main(int argc, char * argv[]) {
286         time_t t1 = time(NULL);
287         Init();
288         ReadInputFile(argv[1]);
289     }
```

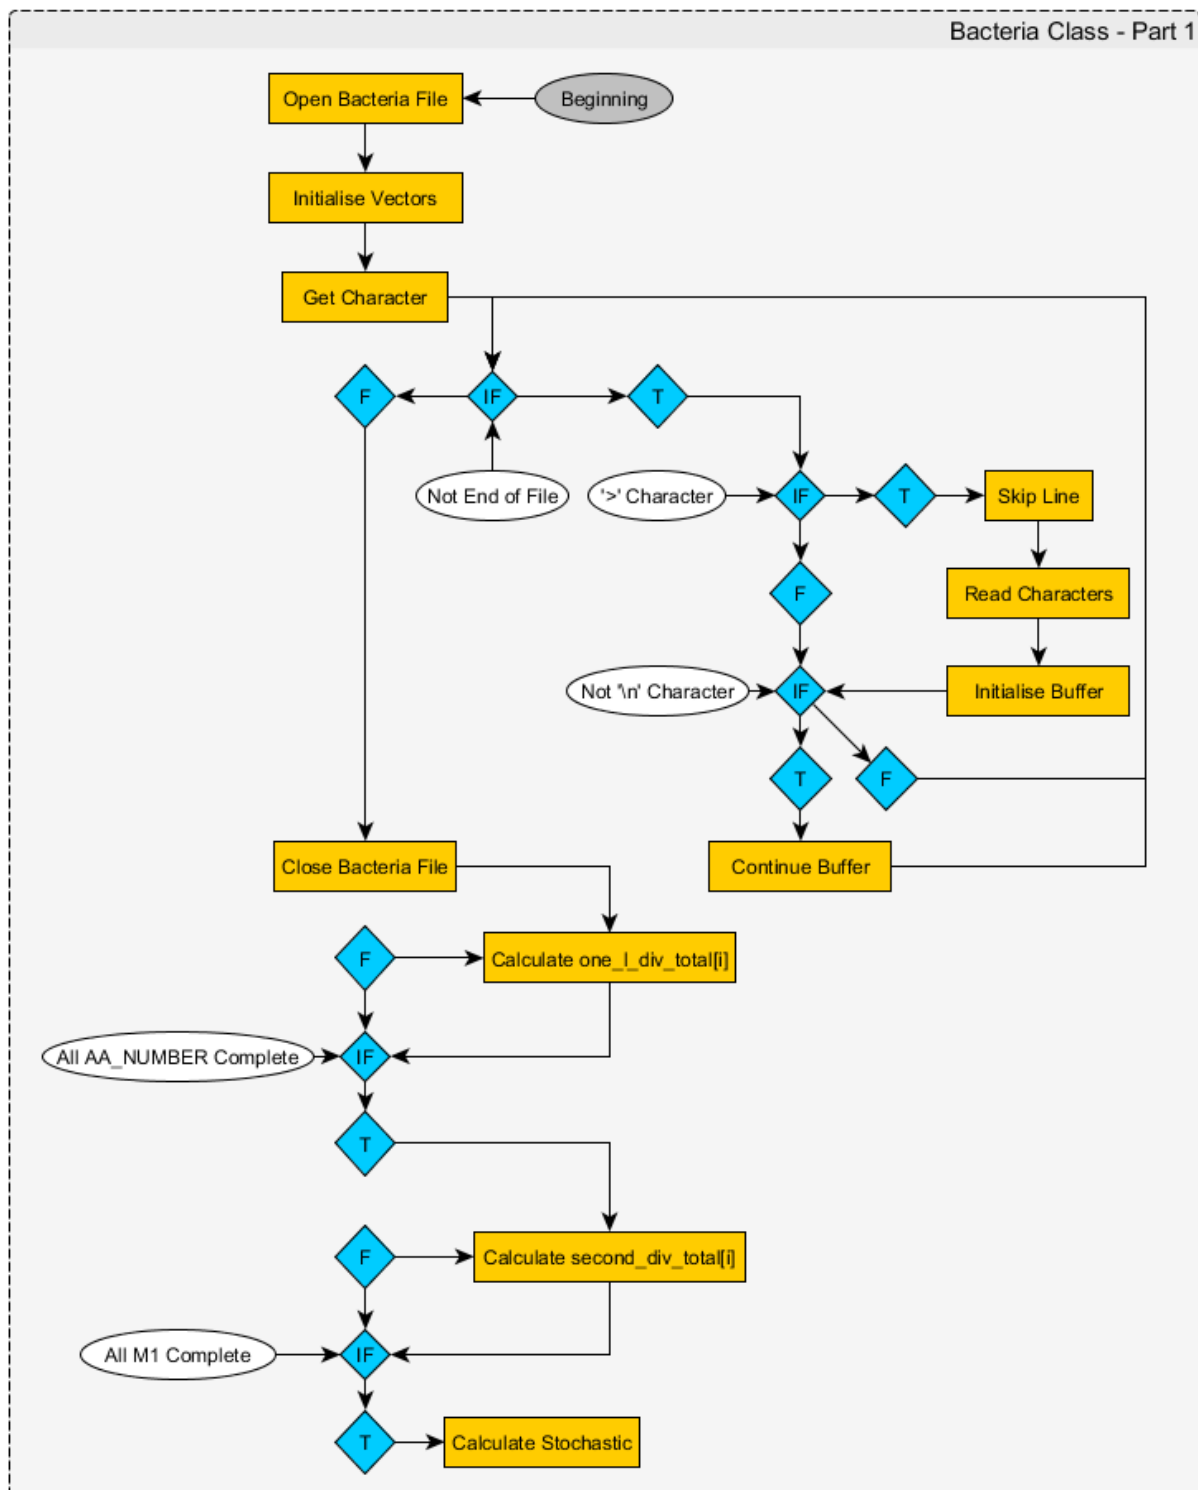


```
366         CompareAllBacteria();
367
368         time_t t2 = time(NULL);
369         printf("time elapsed: %d seconds\n", t2 - t1);
370
371         system("pause");
372
373         return 0;
374     }
```

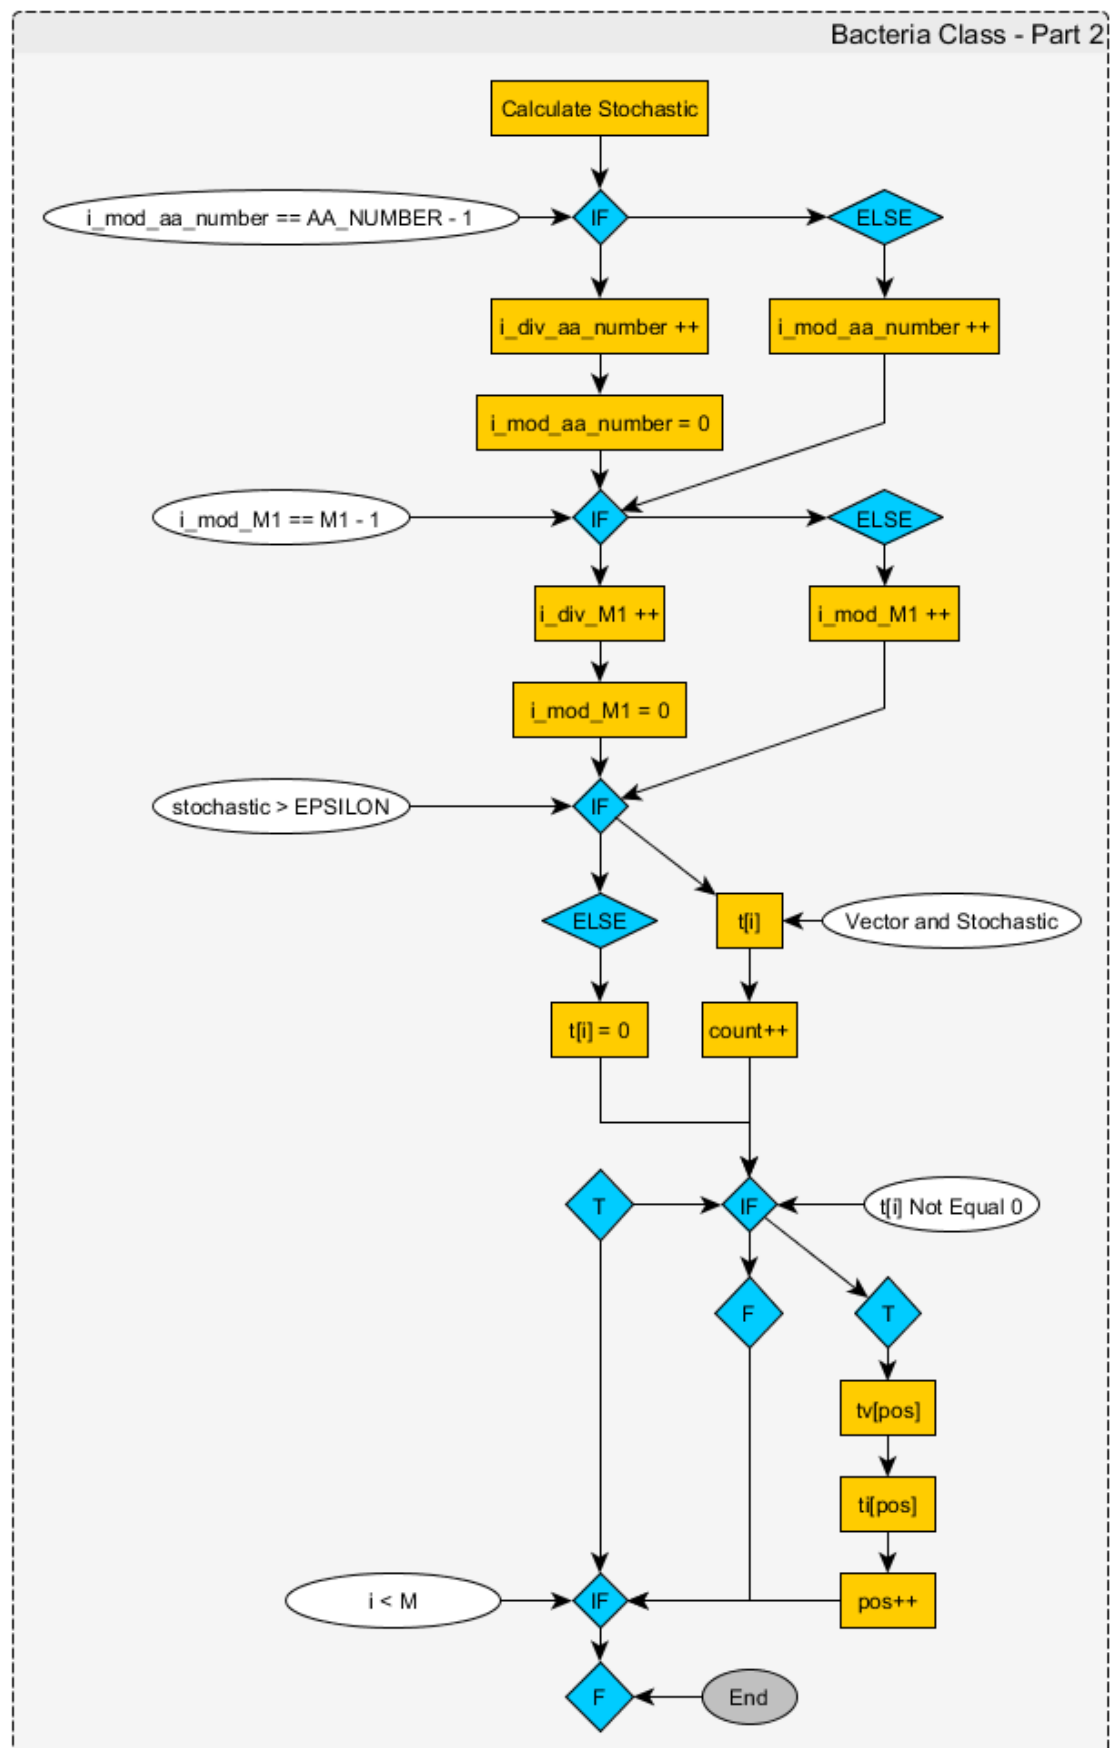
6.3 Detailed Software Architecture



6.4 Bacteria Class – Part 1 Software Architecture



6.5 Bacteria Class – Part 2 Software Architecture



6.6 Compare Bacteria Software Architecture

